

Variational Laplace with Low-Rank and Heteroscedastic Noise Modeling for Nonlinear Dynamical Systems

Dr Alexander Shaw

University of Exeter

Abstract

This manuscript describes and contrasts two variational inference routines used for Dynamic Causal Modelling (DCM): the standard SPM implementation `spm_nlsi_GN.m`, and the extended routines `fitVL_LowRankNoise.m` and `fitVariationalLaplaceThermo.m`. These introduce multiple innovations, including low-rank approximations to both the posterior and observation noise covariances, structured noise smoothing via (Gaussian) radial basis functions, and fallback strategies for numerical stability. These improvements aim to increase robustness, computational efficiency, and biological plausibility in the estimation of hierarchical generative models.

1 Introduction to Variational Laplace

Bayesian inference in nonlinear dynamical systems is challenging due to the intractability of computing exact posterior distributions. Variational Laplace (VL) is a practical and widely used solution to this problem, especially in neuroscience, where it underlies Dynamic Causal Modelling (DCM) and related approaches in the SPM software suite.

VL belongs to the family of variational inference methods, which approximate the true posterior $p(\mathbf{m} \mid \mathbf{y})$ over parameters \mathbf{m} given data \mathbf{y} by a simpler distribution $q(\mathbf{m})$, typically a multivariate Gaussian. The quality of the approximation is measured via the Evidence Lower Bound (ELBO):

$$\mathcal{F} = \mathbb{E}_q[\log p(\mathbf{y}, \mathbf{m})] - \mathbb{E}_q[\log q(\mathbf{m})] = \log p(\mathbf{y}) - \text{KL}(q(\mathbf{m}) \parallel p(\mathbf{m} \mid \mathbf{y})) \quad (1)$$

Maximising \mathcal{F} both approximates the posterior and provides a bound on the marginal likelihood $\log p(\mathbf{y})$, which is useful for model comparison.

1.1 Gaussian Approximation and the Laplace Assumption

In the Laplace approximation, the variational distribution $q(\mathbf{m})$ is assumed to be Gaussian:

$$q(\mathbf{m}) = \mathcal{N}(\mathbf{m} \mid \boldsymbol{\mu}, \boldsymbol{\Sigma}) \quad (2)$$

Here, $\boldsymbol{\mu}$ is the mean of the posterior and $\boldsymbol{\Sigma}$ its covariance. Under this assumption, the ELBO becomes analytically tractable, with three principal contributions:

$$\mathcal{F} = \log p(\mathbf{y} \mid \boldsymbol{\mu}) - \frac{1}{2}(\boldsymbol{\mu} - \mathbf{m}_0)^\top \mathbf{S}_0^{-1}(\boldsymbol{\mu} - \mathbf{m}_0) + \frac{1}{2} \log \det \boldsymbol{\Sigma} \quad (3)$$

This decomposition highlights three elements: (i) data fit via likelihood, (ii) deviation from the prior, and (iii) posterior uncertainty.

1.2 Model Setup and Assumptions

The model assumes that data are generated by a nonlinear function of parameters:

$$\mathbf{y} = f(\mathbf{m}) + \mathbf{e}, \quad \mathbf{e} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma}_{\text{obs}}) \quad (4)$$

and that parameters are drawn from a Gaussian prior:

$$\mathbf{m} \sim \mathcal{N}(\mathbf{m}_0, \mathbf{S}_0) \quad (5)$$

The function f may describe anything from a sigmoid curve to a full nonlinear neural mass model.

The primary goal of VL is to compute a Gaussian approximation to the posterior over \mathbf{m} , and to optionally estimate the model evidence via the ELBO. This is achieved by iteratively updating the mean $\boldsymbol{\mu}$ and precision $\boldsymbol{\Sigma}^{-1}$ using information about the model's Jacobian and residuals.

2 Standard SPM Implementation: `spm_nlsi_GN.m`

The routine `spm_nlsi_GN.m` forms the backbone of parameter inference in classical DCMs. It implements a variational Laplace scheme with a full-rank Gaussian posterior and scalar (homoscedastic) observation noise. The updates proceed in an Expectation-Maximisation-like loop where posterior parameters are updated until convergence.

2.1 Linearisation and Gauss-Newton Approximation

To render the ELBO tractable, the function $f(\mathbf{m})$ is linearised around the current estimate $\boldsymbol{\mu}$ via a first-order Taylor expansion:

$$f(\mathbf{m}) \approx f(\boldsymbol{\mu}) + \mathbf{J}(\mathbf{m} - \boldsymbol{\mu}) \quad (6)$$

where $\mathbf{J} \in \mathbb{R}^{n \times d}$ is the Jacobian of f evaluated at $\boldsymbol{\mu}$. This gives a local linear model that is used to construct a quadratic approximation to the ELBO.

2.2 Posterior Update Steps

Using the linearised model, the posterior precision is given by:

$$\boldsymbol{\Sigma}^{-1} = \mathbf{J}^\top \boldsymbol{\Sigma}_{\text{obs}}^{-1} \mathbf{J} + \mathbf{S}_0^{-1} \quad (7)$$

The posterior mean is then updated as:

$$\boldsymbol{\mu}_{\text{new}} = \boldsymbol{\mu}_{\text{old}} + \boldsymbol{\Sigma} \left(\mathbf{J}^\top \boldsymbol{\Sigma}_{\text{obs}}^{-1} (\mathbf{y} - f(\boldsymbol{\mu})) - \mathbf{S}_0^{-1} (\boldsymbol{\mu} - \mathbf{m}_0) \right) \quad (8)$$

These steps are iterated until convergence of $\boldsymbol{\mu}$ or until a maximum number of iterations is reached.

2.3 Limitations and Computational Bottlenecks

While elegant and effective in many applications, the standard implementation has several limitations:

- **Homoscedastic Noise Assumption:** The routine assumes a scalar noise precision λ , i.e. $\Sigma_{\text{obs}} = \lambda^{-1}\mathbf{I}$. This cannot capture structured or time-varying noise often present in EEG/MEG data.
- **Full-Rank Posterior:** The full posterior covariance matrix must be explicitly computed and stored, scaling as $O(d^2)$ in memory and $O(d^3)$ in computation. For high-dimensional DCMs or hierarchical models, this becomes infeasible.
- **Static Noise Model:** The noise precision is typically updated outside the inference loop using restricted maximum likelihood (ReML). This separates noise estimation from inference and can result in unstable or delayed updates.
- **No Entropy Tracking:** The entropy of the posterior $\frac{1}{2} \log \det \Sigma$ is not explicitly monitored. This makes it difficult to diagnose convergence issues or posterior overconfidence.

These limitations motivate the development of more flexible and scalable routines, such as `fitVL_LowRankNoise`, described next.

3 Extensions in `fitVL_LowRankNoise.m` and `fitVariationalLaplaceThermo.m`

The function `fitVariationalLaplaceThermo.m` extends the classical variational Laplace routine by introducing a more expressive posterior approximation and a dynamic, structured noise model. These enhancements are specifically designed to address known limitations of the SPM routine and to increase robustness when fitting hierarchical, noisy, or high-dimensional models.

3.1 Low-Rank Posterior Covariance

The first major innovation is to replace the full-rank posterior covariance Σ with a structured approximation:

$$\Sigma \approx \mathbf{V}\mathbf{V}^\top + \text{diag}(\mathbf{D}) \quad (9)$$

This formulation draws inspiration from methods in Gaussian process inference, principal component analysis (PCA), and low-rank Bayesian filtering. It is based on the observation that, in many applied models, posterior uncertainty is concentrated in a small number of directions. In a neural model, for instance, these directions may correspond to shared synaptic or circuit-level parameters.

Benefits:

- Reduces storage complexity from $O(d^2)$ to $O(dk)$, where $k \ll d$
- Enables efficient inversion and entropy estimation using the Woodbury identity
- Supports interpretable dimensionality reduction of uncertainty

Posterior Update: Given the precision matrix \mathbf{H} from the ELBO Hessian, the low-rank posterior is formed via SVD:

$$\mathbf{H} \approx \mathbf{U}_k \mathbf{\Lambda}_k \mathbf{U}_k^\top, \quad \mathbf{V} = \mathbf{U}_k \mathbf{\Lambda}_k^{-1/2}$$

The diagonal correction \mathbf{D} is set such that the marginal variances are preserved:

$$D_i = H_{ii} - \sum_j V_{ij}^2$$

This ensures numerical consistency and interpretable uncertainty quantification.

3.2 Structured Observation Noise: Low-Rank + Diagonal

Traditional VL assumes independent and identically distributed (i.i.d.) Gaussian noise, where:

$$\mathbf{\Sigma}_{\text{obs}} = \sigma^2 \mathbf{I}$$

This assumption fails in the presence of autocorrelated noise, physiological artifacts (e.g. eye blinks, cardiac rhythms), or temporally varying signal-to-noise ratios.

To address this, in `fitVL.LowRankNoise.m`, we further introduce a structured noise model:

$$\mathbf{\Sigma}_{\text{obs}} \approx \mathbf{U} \mathbf{U}^\top + \text{diag}(\mathbf{D}_{\text{noise}}) \quad (10)$$

Here, $\mathbf{U} \in \mathbb{R}^{n \times k}$ represents low-dimensional latent noise processes, and $\mathbf{D}_{\text{noise}} \in \mathbb{R}^n$ captures heteroscedastic per-observation variance. This hybrid model enables both global and local noise structure to be inferred simultaneously.

Adaptation Strategy: At each iteration:

1. Residuals $\mathbf{r} = \mathbf{y} - f(\mathbf{m})$ are computed
2. A smooth covariance kernel is constructed using the residuals (see next section)
3. The top k components of this kernel are extracted as \mathbf{U}
4. The diagonal component $\mathbf{D}_{\text{noise}}$ is updated using a moving average:

$$D^{(t+1)} = \alpha D^{(t)} + (1 - \alpha) \max(r^2 - \|\mathbf{U}_i\|^2, \epsilon)$$

This approach enables structured learning of noise covariance that is responsive to model fit quality.

3.3 Noise Smoothing via Radial Basis Functions

To prevent overfitting noise to isolated outliers and to regularise the evolving noise model, we smooth residuals using radial basis functions (RBFs). Let the smoothing kernel be:

$$K_{ij} = \exp\left(-\frac{(x_i - x_j)^2}{2\ell^2}\right) \quad (11)$$

where x_i are observation indices or timestamps and ℓ is a length-scale hyperparameter (typically fixed at 1–3).

The smoothed residual kernel is then:

$$\mathbf{K} \cdot \text{diag}(\mathbf{r}) \cdot \mathbf{K}^\top$$

From this kernel, we extract \mathbf{U} via SVD or eigendecomposition. This gives a low-rank noise structure aligned with locally coherent deviations from the model.

Rationale:

- Smooths the residual energy landscape
- Encourages structured (non-i.i.d.) noise estimation
- Reduces the risk of rank-deficient noise covariance estimates

In practice, this step stabilises convergence and improves robustness across a wide range of nonlinear models.

4 Numerical Stability Enhancements

Many practical implementations of VL suffer from breakdowns due to ill-conditioning of the observation noise covariance matrix Σ_{obs} , especially early in inference or with poor priors. We address this with two key techniques: jittered Cholesky and fallback Woodbury inversion.

4.1 Jittered Cholesky Decomposition

Cholesky decomposition of Σ_{obs} is used to compute the inverse and determinant required for ELBO evaluation. However, Σ_{obs} may not be positive definite due to numerical noise or unbalanced updates.

We therefore apply:

$$\mathbf{L} = \text{chol}(\Sigma_{\text{obs}} + \epsilon \mathbf{I}) \tag{12}$$

If Cholesky fails, ϵ is increased iteratively (e.g. $1e-6 \rightarrow 1e-4 \rightarrow 1e-2$) until success is achieved.

4.2 Woodbury Identity Fallback

If Cholesky still fails after multiple attempts, we switch to a robust inversion using the Woodbury identity:

$$(\mathbf{U}\mathbf{U}^\top + \mathbf{D})^{-1} = \mathbf{D}^{-1} - \mathbf{D}^{-1}\mathbf{U} \left(\mathbf{I} + \mathbf{U}^\top \mathbf{D}^{-1} \mathbf{U} \right)^{-1} \mathbf{U}^\top \mathbf{D}^{-1} \tag{13}$$

This provides numerically stable inversion even when $\mathbf{U}\mathbf{U}^\top$ is nearly low-rank or \mathbf{D} has near-zero entries.

Advantages:

- Avoids full matrix inversion ($O(n^3)$)
- Enables stability in early inference iterations
- Ensures meaningful ELBO and gradient evaluations even under degeneracy

5 Thermodynamic Free Energy Tracking

In `fitVL_LowRankNoise`, the Evidence Lower Bound (ELBO) is tracked explicitly at each iteration. This improves transparency, convergence diagnostics, and future integration with annealing or thermodynamic integration schemes.

5.1 ELBO Decomposition

The total ELBO is given by:

$$\mathcal{F} = \underbrace{\log p(\mathbf{y} \mid \mathbf{m})}_{\text{likelihood}} + \underbrace{\log p(\mathbf{m})}_{\text{prior}} + \underbrace{H[q(\mathbf{m})]}_{\text{entropy}} \quad (14)$$

Each term is computed as:

Likelihood:

$$\log p(\mathbf{y} \mid \mathbf{m}) = -\frac{1}{2}(\mathbf{y} - f(\mathbf{m}))^\top \boldsymbol{\Sigma}_{\text{obs}}^{-1}(\mathbf{y} - f(\mathbf{m})) - \frac{1}{2} \log \det \boldsymbol{\Sigma}_{\text{obs}} - \frac{n}{2} \log 2\pi$$

Prior:

$$\log p(\mathbf{m}) = -\frac{1}{2}(\mathbf{m} - \mathbf{m}_0)^\top \mathbf{S}_0^{-1}(\mathbf{m} - \mathbf{m}_0)$$

Entropy (low-rank):

$$H[q(\mathbf{m})] = \frac{1}{2} \log \det(\mathbf{V}\mathbf{V}^\top + \text{diag}(\mathbf{D})) + \frac{d}{2}(1 + \log 2\pi)$$

These values are stored across iterations, allowing for detailed inspection of convergence behavior.

Thermodynamic Use: The method also supports free-energy annealing by introducing a temperature parameter β , which controls the sharpness of the likelihood term. This provides a natural mechanism for thermodynamic integration or variational tempering in future extensions.

6 Pseudocode for `fitVL_LowRankNoise`

The following pseudocode summarises the algorithm:

Input: Observed data \mathbf{y} , model $f(\mathbf{m})$, prior mean \mathbf{m}_0 , prior covariance \mathbf{S}_0

Output: Posterior mean \mathbf{m} , low-rank posterior factors \mathbf{V} , \mathbf{D} , log-likelihood $\log L$

1. Initialize $\mathbf{m} = \mathbf{m}_0$
2. Initialize posterior structure:
 - \mathbf{V}_{post} from truncated SVD of \mathbf{S}_0
 - $\mathbf{D}_{\text{post}} = \text{diag}(\mathbf{S}_0 - \mathbf{V}_{\text{post}} * \mathbf{V}_{\text{post}}^\top)$
3. Initialize noise model:
 - Compute residuals $\mathbf{r} = \mathbf{y} - f(\mathbf{m})$
 - Use RBF kernel to build smoothed covariance matrix
 - Extract low-rank noise factors $\mathbf{U}_{\text{noise}}$ via SVD

```

    Set D_noise = initial diagonal noise estimate

4. For iter = 1 to maxIter:
    a. Predict y_pred = f(m)
    b. Compute residuals = y - y_pred
    c. Build noise covariance: Sigma = U_noise * U_noise^T + diag(D_noise)
    d. Try Cholesky decomposition:
        If fails, fallback to Woodbury identity
    e. Compute Jacobian J = df/dm
    f. Compute precision H = J^T * inv(Sigma) * J
    g. Compute gradient g = J^T * inv(Sigma) * residuals - prior gradient
    h. Update mean m ← m + inv(H) * g
    i. Update posterior covariance:
        V_post ← SVD(H)
        D_post ← diag(H - V_post * V_post^T)
    j. Update D_noise and U_noise using residuals + RBF smoothing
    k. Track ELBO and check ||dm|| < tol

5. Return final m, V_post, D_post, Sigma_struct, ELBO trace

```

7 Synthetic Data Example

To illustrate the benefits of the method, we simulate data from a simple nonlinear model with heteroscedastic noise:

$$y_i = \sin(mx_i) + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma_i^2)$$

where:

- $m = 2$ (true frequency)
- $x_i \in [0, 2\pi]$, sampled at 100 points
- $\sigma_i = 0.05 + 0.2 \cdot \frac{i}{n}$ (linearly increasing noise)
- Prior: $m \sim \mathcal{N}(1, 0.5^2)$

Results

The algorithm:

- Recovered the posterior mean $\hat{m} = 2.01$
- Correctly modeled the increasing observation variance
- Showed monotonic improvement in ELBO
- Output interpretable low-rank noise structure \mathbf{U}

This simple case highlights the power of the algorithm to model both latent structure and adaptive noise within a unified variational framework.

8 Conclusion

We presented a structured variational inference routine, `fitVL_LowRankNoise`, that extends classical variational Laplace by introducing low-rank approximations for both the posterior and observation noise covariances. We showed how this method overcomes limitations of existing implementations by increasing scalability, improving robustness, and capturing heteroscedastic and structured noise patterns common in biological data.

This method is particularly useful for hierarchical models, dynamic systems with nonstationary noise, and inference tasks requiring principled uncertainty quantification. Future directions include integration with variational annealing, extensions to mixture posteriors, and full Bayesian hierarchical inference.